

Investigating Instruction Pipelines

Introduction

Objectives

At the end of this lab you should be able to:

- Demonstrate the difference between pipelined and sequential processing of the CPU instructions
- Explain pipeline data dependency and data hazard
- Describe a pipeline technique to eliminate data hazards
- Demonstrate compiler “loop unrolling” optimization’s benefits for instruction pipelining
- Describe compiler re-arranging instructions to minimize data dependencies
- Show the use of jump-predict table for pipeline optimisation

Basic Theory

Modern CPUs incorporate instruction pipelines which are able to process different stages of multiple-stage instructions in parallel thus improving the overall performance of the CPUs. However, most programs include instructions that do not readily lend themselves to smooth pipelining thus causing pipeline hazards and effectively reducing the CPU performance. As a result CPU pipelines are designed with some tricks up their sleeves for dealing with these hazards.

Lab Exercises - Investigate and Explore

The lab investigations are a series of exercises that are designed to demonstrate the various aspects of CPU instruction pipelining.

Exercise 1 – Difference between the sequential and the pipelined execution of CPU instructions

Enter the following source code, compile it and load in simulator's memory:

```
program Ex1
  for n = 1 to 20
    p = p + 1
  next
end
```

Open the CPU pipeline window by clicking on the **SHOW PIPELINE...** button in the CPU simulator's window. You should now see the **Instruction Pipeline** window. This window simulates the behaviour of a CPU pipeline. Here we can observe the different stages of the pipeline as program instructions are processed. This pipeline has five stages. The stages are colour-coded as shown in the key for the "Pipeline Stages".

List the names of the stages here:

Fetch	Execute
Decode	Write Result
Read Operands	

The instructions that are being pipelined are listed on the left side (in white text boxes). The newest instruction in the pipeline is at the bottom and the oldest at the top. You'll see this when you run the instructions. The horizontal yellowish boxes display the stages of an instruction as it progresses through the pipeline. At the bottom left corner pipeline statistics are displayed as the instructions are executed.

Check the box titled **Stay on top** and make sure **No instruction pipeline** check box is selected. In the CPU simulator window bring the speed slider down to around a reading of 30. Run the program and observe the pipeline. Wait for the program to complete. Now make a note of the following values

CPI (Clocks Per Instruction)	5.04
SF (Speed-up Factor)	0.99

Next, uncheck the **No instruction pipeline** checkbox, reset and run the above program again and wait for it to complete.

Note down your observation on how the pipeline visually behaved differently

The next instruction already starts executing before the previous one has finished. If executing a stage is not possible yet, this is indicated by bubbles, hazards and syncs.

Now once again make a note of the following values

CPI (Clocks Per Instruction)	1.92
SF (Speed-up Factor)	2.6

Briefly explain why you think there is a difference in the two sets of values:

The instructions are pipelined, so the stages of the pipeline do not have to wait for each other.

Exercise 2 – CPU pipeline data hazards, bubbles and the NOP instruction

CPU pipelines often have to deal with various hazards, i.e. those aspects of the CPU architecture which prevent the pipelines running smoothly and uninterrupted. These are often called “pipeline hazards”. One such hazard is called the “data hazard”. A data hazard is caused by unavailability of an operand value when it is needed. In order to demonstrate this create a program (call it Ex2) and enter the following set of instructions

```
MOV #1, R01
MOV #5, R03
MOV #3, R01
ADD R01, R03
HLT
```

Before you carry on, make a note of what value you expect to see in register R03 at the end of the above set of instructions. Make a note of it below:

R03 = 8

Make sure the **No instruction pipeline** is NOT checked and **Do not insert bubbles** is checked. Reset the program and run the above instructions. Make a note of the value in register R03 below:

R03 = 6

Now insert a NOP instruction (use the **Miscellaneous** tab) after the third instruction, i.e. you end up with the following modified set of instructions

```
MOV #1, R01
MOV #5, R03
MOV #3, R01
NOP
ADD R01, R03
HLT
```

Reset the program and run the above set of instructions. Observe the value in register R03 when the program completes. Make a note of this value below

R03 = 8

You have now recorded three instances of the values of the register R03 the first one being your prediction. Briefly comment on your observations of the three values:

In the second one there was a data hazard. Three was not yet written back to register 1 when the add operation read its input.

Now delete the NOP instruction from above program and uncheck the option **Do not insert bubbles**. Reset the program and run the instructions. Observe the value in register R03 when the program completes. Make a note of this value below:

R03 = 8

The value above should be the same as the one when you inserted a NOP instruction in the program. However this value is obtained without the NOP instruction in this case. Briefly explain:

The CPU automatically detects data hazards and stalls the pipeline.

Have you seen the “bubble”? What colour is it?

Red

Finally, make a note of the following values:

CPI (Clocks Per Instruction)	2.6
SF (Speed-up Factor)	1.92
Data Hazards	1

Exercise 3 – A pipeline technique to eliminate data hazards

One way of dealing with “data hazard” is to get the CPU to “speed up” the availability of operands to pipelined instructions. One such method is called “operand forwarding”, a kind of short-cut by the hardware. To demonstrate this check the box titled **Enable operand forwarding** and run the above code again.

Has the bubble in Exercise 4 disappeared or burst?

Disappeared

The simulator keeps a count of the pipeline hazards it detects as the instructions go through the pipeline. These can be seen near the bottom of the pipeline window.

Make a note of the following values

CPI (Clocks Per Instruction)	2.6
SF (Speed-up Factor)	1.92
Data Hazards	0

Has there been an improvement?

There is no improvement.

Exercise 4 – Loop unrolling optimization minimizing control dependencies

In a previous tutorial on compiler optimizations, we looked at one method of optimization called “loop unrolling”. This method essentially duplicates the inner code of a loop as many times as the number of loops, removing some redundant code as well as the loop’s compare and jump instructions. However, the code size of the program increases. It is shown that “loop unrolling” is well suited to instruction pipelining which takes full advantage of it thus improving CPU performance. Here, we will prove this to be the case.

Enter the following code, select optimization option **Redundant Code** and compile it.

```
program Ex4_1
  for n = 1 to 8
    t = t + 1
  next
end
```

Make a note of the size of the code generated for Ex4_1 here:

32

Now, load this code in CPU simulator's memory.

Next, make sure the optimization option **Loop Unrolling** is selected in addition to the option **Redundant Code** optimization. Change the program name to Ex4_2 and compile it again. Load this code in memory too. So, now you should have two versions of the code: Ex4_1 without "loop unrolling" optimization and Ex4_2 with "loop unrolling" optimization.

Make a note of the size of the code generated for Ex4_2 here:

109

Make sure the pipeline window stays on top. Also make sure the **Enable operand forwarding** and **Enable jump prediction** boxes are all unchecked. First, select program Ex4_1 from the **PROGRAM LIST** frame in the CPU simulator window then click the **RESET** button. Make sure the speed of simulation is set at maximum. Now click the **RUN** button to run program Ex4_1. Observe the pipeline and when the program is finished make a note of the following values:

CPI (Clocks Per Instruction)	5.09
SF (Speed-up Factor)	0.98
No of instructions executed	44

Do the same with program Ex4_2 and make note of the following values:

CPI (Clocks Per Instruction)	5.63
SF (Speed-up Factor)	0.89
No of instructions executed	19

Briefly comment on your observations making references to the code sizes and the number of instructions executed:

The loop unrolling resulted in a bigger code size but the amount of instructions executed is decreased. This is because the code needs less jump and compare instructions.

Exercise 5 – Compiler re-arranging instruction sequence to help minimize data dependencies

The optimization in Exercise 4 is one example of how a modern compiler can provide support for the CPU pipeline. Another example is when the compiler re-arranges the code without changing the logic of the code. This is done to minimize pipeline hazards such as the “data hazard” we studied in Exercise 3. Here we demonstrate this technique.

Make sure **Show dependencies** check box is checked and ONLY the **Redundant Code** optimization is selected. Enter the following source code, compile it and load in memory

```
program Ex5_1
    a = 1
    b = a
    c = 2
end
```

Copy the CPU instruction sequence generated below (do not include the instruction addresses):

```
MOV #1, R01
MOV R01, R02
MOV #2, R03
HLT
```

Next, select the optimization option **Code Dependencies**. Change the program name to Ex5_2, compile it and load in memory.

Copy the CPU instruction sequence generated below:

```
MOV #1, R01
MOV #2, R03
MOV R01, R02
HLT
```

How do the two sequences differ? Does the change affect the logic of the program? Briefly explain the rationale for the change:

The order differs. By moving the middle instruction to the back the data hazard for R01 is eliminated.

Let's see if we can measure the improvement with this "out of sequence execution" method applied. Modify the above program as below:

```
program Ex5_3
  for n = 1 to 50
    a = 1
    b = a
    c = 2
  next
end
```

Now, compile and load two version of the above program, one without the **Code Dependencies** optimization and one with this optimization (call this one program Ex5_4).

First run program Ex5_3 and make note of the values below:

CPI (Clocks Per Instruction)	1.74
SF (Speed-up Factor)	2.87

Next, run program Ex5_4 and make note of the values below:

CPI (Clocks Per Instruction)	1.6
SF (Speed-up Factor)	3.12

Do you see any improvement in program Ex5_4 over program Ex5_3 (express this in percentage)?

The speed up is about 9%. This is because the bubbles are not needed.

Exercise 6 – Jump predict table

The CPU pipeline uses a table to keep a record of the predicted jump addresses. So, whenever a conditional jump instruction is being executed this table is consulted in order to see what the jump address is predicted as. If this prediction is wrong then the calculated address is used instead. Often the predicted address will be correct with occasional wrong prediction. However, the overall effect will be an improvement on CPU's performance.

Enter the following program and compile it with ONLY the **Enable optimizer** and **Remove redundant code** check boxes selected. Load the compiled program in the CPU.

```
program Ex6
  i = 0
  for p = 1 to 40
    i = i + 1
    if i = 10 then
      i = 0
      r = i
    end if
  next
end
```

Run the program and make a note of the following pipeline stats:

CPI (Clocks Per Instruction)	2.21
SF (Speed-up Factor)	2.26

Now, in the pipeline window select the **Enable jump prediction** check box. Reset the program and run it again. Make a note of the following pipeline stats:

CPI (Clocks Per Instruction)	1.95
SF (Speed-up Factor)	2.56

Do you see a difference? Is it an improvement?

Yes, there is an improvement.

Click on the **SHOW JUMP TABLE...** button. You should see the **Jump Predict Table** window showing. This table keeps an entry relevant to each conditional jump instruction. The information contained has the following fields. Can you suggest what each field stands for? Enter your suggestions in the table below:

V	'0' if the jump was not taken in the last 2 executions of the jump instruction, else '1'.
JInstAddr	Address where jump instruction is located in the program code.
JTarget	Target address of the jump instruction.
PStat	V='0': PStat is also '0'. V='1': PStat increases when jump is taken (max '2') and decreases when jump is not taken.
Count	Total amount of times the jump was taken. Reset by V='0'.

Note: The CPU only predicts the jump is taken (i.e.: fill the pipeline with the instructions from the jump target address and further) when PStat='2'!